

Convolutional Neural Network with Numpy (Fast)



CNNumpy

In the [previous post](https://hackmd.io/@machine-learning/blog-post-cnnumpy-slow) (<https://hackmd.io/@machine-learning/blog-post-cnnumpy-slow>), we have seen a naive implementation of Convolutional Neural network using Numpy.

Here, we are going to implement a faster CNN using Numpy with the im2col/col2im method.

To see the full implementation, please refer to my [repository](https://github.com/3outeille/CNNumpy) (<https://github.com/3outeille/CNNumpy>).

Also, if you want to read some of my blog posts, feel free to check them at my [blog](https://3outeille.github.io/deep-learning/) (<https://3outeille.github.io/deep-learning/>).

I) Forward propagation

- The main objective here is to ensure that the reader develops a strong intuition about how im2col/col2im works so that he can write them by himself.
- Thus, I decided to be less **rigorous** in my explanation to make things **clearer**.
- We will refer to the term “**level**” as a whole horizontal kernel slide from left to right.
- We will only discuss **average pooling** layer here (even though same logic can be applied on max pooling).

1) Convolutional layer

- In the naive implementation, we used a lot of nested “For loops” which makes our code very slow.
- An approach could be to trade some memory for more speedup.
- Here are the following steps:
 - **A.** Transform our input image into a matrix (im2col).
 - **B.** Reshape our kernel (flatten).
 - **C.** Perform matrix multiplication between reshaped input image and kernel.
- We are going to see how it works intuitively and then how to implement it using Numpy.

≡ Intuition

- As an example, we will perform a convolution between an (1,3,4,4) input image and kernels of shape (2,3,2,2).

A) Transform our input image into a matrix (im2col)

- Here is how it works:

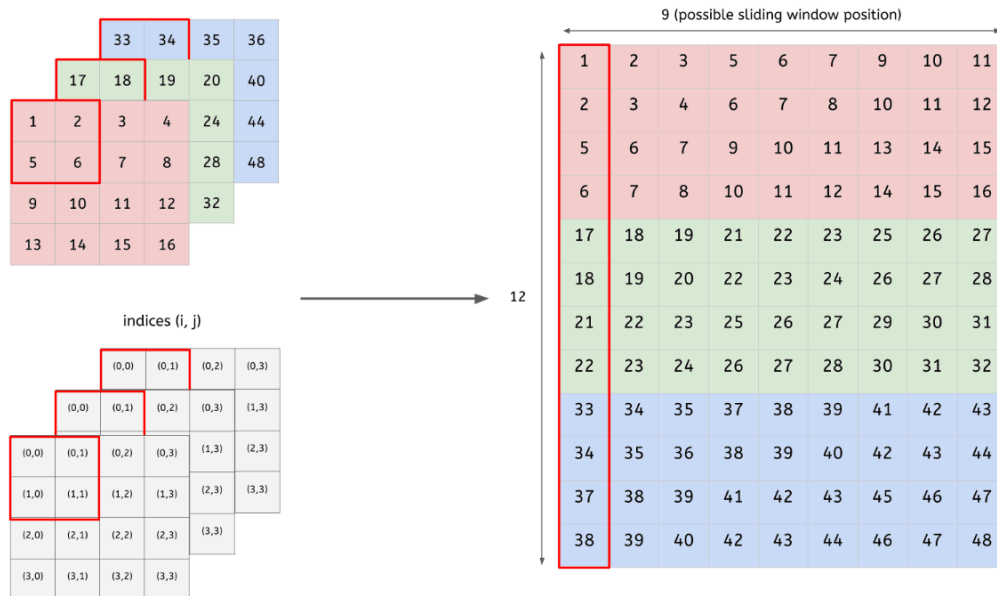



Figure 1: Input image transformed into matrix

- How do we do that in Numpy? An efficient way to do so is by the help of multi-dimensional arrays indexing (<https://docs.scipy.org/doc/numpy/user/basics.indexing.html>). For example,

```
>>> y = np.arange(35).reshape(5,7)
>>> y
array([[ 0,  1,  2,  3,  4,  5,  6],
       [ 7,  8,  9, 10, 11, 12, 13],
       [14, 15, 16, 17, 18, 19, 20],
       [21, 22, 23, 24, 25, 26, 27],
       [28, 29, 30, 31, 32, 33, 34]])
>>> y[np.array([0,2,4]), np.array([0,1,2])]
array([0, 15, 30]) # 0 = (0,0) / 15 = (2,1) / 30 = (4,2)
```

- Thus, the idea is to use **the multi-dimensional arrays indexing** property to transform our input image into a matrix.
- Indeed, we can notice few things:
 - Firstly**, the indices for each input image channel is the same. Thus, we can focus ourselves only on the first channel since the result will be the same for the others.



◦ **Secondly**, we can observe a certain pattern in the indices (i, j) when we slide our kernel.

■ Index i:

■ At level 1, we have:

indices (i, j)

(0,0)	(0,1)	(0,2)	(0,3)
(1,0)	(1,1)	(1,2)	(1,3)
(2,0)	(2,1)	(2,2)	(2,3)
(3,0)	(3,1)	(3,2)	(3,3)

	j									
	(0,0)	(0,1)	(0,2)	(1,0)	(1,1)	(1,2)	(2,0)	(2,1)	(2,2)	
	(0,1)	(0,2)	(0,3)	(1,1)	(1,2)	(1,3)	(2,1)	(2,2)	(2,3)	
	(1,0)	(1,1)	(1,2)	(2,0)	(2,1)	(2,2)	(3,0)	(3,1)	(3,2)	
	(1,1)	(1,2)	(1,3)	(2,1)	(2,2)	(2,3)	(3,1)	(3,2)	(3,3)	
	(0,0)	(0,1)	(0,2)	(1,0)	(1,1)	(1,2)	(2,0)	(2,1)	(2,2)	
	(0,1)	(0,2)	(0,3)	(1,1)	(1,2)	(1,3)	(2,1)	(2,2)	(2,3)	
	(1,0)	(1,1)	(1,2)	(2,0)	(2,1)	(2,2)	(3,0)	(3,1)	(3,2)	
	(1,1)	(1,2)	(1,3)	(2,1)	(2,2)	(2,3)	(3,1)	(3,2)	(3,3)	
	(0,0)	(0,1)	(0,2)	(1,0)	(1,1)	(1,2)	(2,0)	(2,1)	(2,2)	
	(0,1)	(0,2)	(0,3)	(1,1)	(1,2)	(1,3)	(2,1)	(2,2)	(2,3)	
	(1,0)	(1,1)	(1,2)	(2,0)	(2,1)	(2,2)	(3,0)	(3,1)	(3,2)	
	(1,1)	(1,2)	(1,3)	(2,1)	(2,2)	(2,3)	(3,1)	(3,2)	(3,3)	


■ At level 2, we have:

indices (i, j)



(0,0)	(0,1)	(0,2)	(0,3)
(1,0)	(1,1)	(1,2)	(1,3)
(2,0)	(2,1)	(2,2)	(2,3)
(3,0)	(3,1)	(3,2)	(3,3)

	j									
	(0,0)	(0,1)	(0,2)	(1,0)	(1,1)	(1,2)	(2,0)	(2,1)	(2,2)	
	(0,1)	(0,2)	(0,3)	(1,1)	(1,2)	(1,3)	(2,1)	(2,2)	(2,3)	
	(1,0)	(1,1)	(1,2)	(2,0)	(2,1)	(2,2)	(3,0)	(3,1)	(3,2)	
	(1,1)	(1,2)	(1,3)	(2,1)	(2,2)	(2,3)	(3,1)	(3,2)	(3,3)	
	(0,0)	(0,1)	(0,2)	(1,0)	(1,1)	(1,2)	(2,0)	(2,1)	(2,2)	
	(0,1)	(0,2)	(0,3)	(1,1)	(1,2)	(1,3)	(2,1)	(2,2)	(2,3)	
	(1,0)	(1,1)	(1,2)	(2,0)	(2,1)	(2,2)	(3,0)	(3,1)	(3,2)	
	(1,1)	(1,2)	(1,3)	(2,1)	(2,2)	(2,3)	(3,1)	(3,2)	(3,3)	
	(0,0)	(0,1)	(0,2)	(1,0)	(1,1)	(1,2)	(2,0)	(2,1)	(2,2)	
	(0,1)	(0,2)	(0,3)	(1,1)	(1,2)	(1,3)	(2,1)	(2,2)	(2,3)	
	(1,0)	(1,1)	(1,2)	(2,0)	(2,1)	(2,2)	(3,0)	(3,1)	(3,2)	
	(1,1)	(1,2)	(1,3)	(2,1)	(2,2)	(2,3)	(3,1)	(3,2)	(3,3)	


- indices (i, j)
- | | | | |
|-------|-------|-------|-------|
| (0,0) | (0,1) | (0,2) | (0,3) |
| (1,0) | (1,1) | (1,2) | (1,3) |
| (2,0) | (2,1) | (2,2) | (2,3) |
| (3,0) | (3,1) | (3,2) | (3,3) |



- Index j :
 - At level 1, we have:



- indices (i, j)
- | | | | |
|-------|-------|-------|-------|
| (0,0) | (0,1) | (0,2) | (0,3) |
| (1,0) | (1,1) | (1,2) | (1,3) |
| (2,0) | (2,1) | (2,2) | (2,3) |
| (3,0) | (3,1) | (3,2) | (3,3) |



- At level 3, we have:

indices (i, j)

(0,0)	(0,1)	(0,2)	(0,3)
(1,0)	(1,1)	(1,2)	(1,3)
(2,0)	(2,1)	(2,2)	(2,3)
(3,0)	(3,1)	(3,2)	(3,3)

	j											
i	(0, 0)	(0, 1)	(0, 2)	(1, 0)	(1, 1)	(1, 2)	(2, 0)	(2, 1)	(2, 2)			
	(0, 1)	(0, 2)	(0, 3)	(1, 1)	(1, 2)	(1, 3)	(2, 1)	(2, 2)	(2, 3)			
	(1, 0)	(1, 1)	(1, 2)	(2, 0)	(2, 1)	(2, 2)	(3, 0)	(3, 1)	(3, 2)			
	(1, 1)	(1, 2)	(1, 3)	(2, 1)	(2, 2)	(2, 3)	(3, 1)	(3, 2)	(3, 3)			
	(0, 0)	(0, 1)	(0, 2)	(1, 0)	(1, 1)	(1, 2)	(2, 0)	(2, 1)	(2, 2)			
	(0, 1)	(0, 2)	(0, 3)	(1, 1)	(1, 2)	(1, 3)	(2, 1)	(2, 2)	(2, 3)			
	(1, 0)	(1, 1)	(1, 2)	(2, 0)	(2, 1)	(2, 2)	(3, 0)	(3, 1)	(3, 2)			
	(1, 1)	(1, 2)	(1, 3)	(2, 1)	(2, 2)	(2, 3)	(3, 1)	(3, 2)	(3, 3)			
	(0, 0)	(0, 1)	(0, 2)	(1, 0)	(1, 1)	(1, 2)	(2, 0)	(2, 1)	(2, 2)			
	(0, 1)	(0, 2)	(0, 3)	(1, 1)	(1, 2)	(1, 3)	(2, 1)	(2, 2)	(2, 3)			
	(1, 0)	(1, 1)	(1, 2)	(2, 0)	(2, 1)	(2, 2)	(3, 0)	(3, 1)	(3, 2)			
	(1, 1)	(1, 2)	(1, 3)	(2, 1)	(2, 2)	(2, 3)	(3, 1)	(3, 2)	(3, 3)			
	(0, 0)	(0, 1)	(0, 2)	(1, 0)	(1, 1)	(1, 2)	(2, 0)	(2, 1)	(2, 2)			
	(0, 1)	(0, 2)	(0, 3)	(1, 1)	(1, 2)	(1, 3)	(2, 1)	(2, 2)	(2, 3)			
	(1, 0)	(1, 1)	(1, 2)	(2, 0)	(2, 1)	(2, 2)	(3, 0)	(3, 1)	(3, 2)			
	(1, 1)	(1, 2)	(1, 3)	(2, 1)	(2, 2)	(2, 3)	(3, 1)	(3, 2)	(3, 3)			

■ Conclusion:

- At the level 1, there is a total of 3 slides.
 - For slide 1, we have a [0,1,0,1] vector.
 - For slide 2, we have a [1,2,1,2] vector.
 - For slide 3, we have a [2,3,2,3] vector.
 - **We can notice an increase of 1 at each slide.**
- **At each level, we keep the same pattern.**

- Thus, even if it's not rigorous, we can intuitively think of a general formula for an **(n,n) image convolve to X filters of shape (k, k).**

○ For index i:

- We start with at level 1 with the following vector:

$$\underbrace{[0, 0, \dots, 0, 1, 1, \dots, 1]}_k, \dots, \underbrace{[k-1, k-1, \dots, k-1]}_k$$

- **At each level, we increase this vector by 1**

○ For index j:

- At level 1, there is a total of n-k slides.

- For slide 1, we have a $\underbrace{[0, 1, \dots, k-1, \dots, 0, 1, \dots, k-1]}_k$ vector.

- For slide 2, we have a $\underbrace{[1, 2, \dots, k, \dots, 1, 2, \dots, k]}_k$ vector.

- ...

- For slide n-k, we have a $\underbrace{[n-k, n-k+1, \dots, n-1, \dots, n-k, n-k+1, \dots, n-1]}_k$ vector.

- **At each level, we keep the same pattern.**

- The numbers of filters X do not have any effect in this part. We will see that it's quite simple to deal with it during the kernel reshaping step.
- If you have M images (M > 1), you will have the same matrix but stack horizontally M times.

M = 1

	9										
12	1	2	3	5	6	7	9	10	11		
	2	3	4	6	7	8	10	11	12		
	5	6	7	9	10	11	13	14	15		
	6	7	8	10	11	12	14	15	16		
	17	18	19	21	22	23	25	26	27		
	18	19	20	22	23	24	26	27	28		
	21	22	23	25	26	27	29	30	31		
	22	23	24	26	27	28	30	31	32		
	33	34	35	37	38	39	41	42	43		
	34	35	36	38	39	40	42	43	44		
	37	38	39	41	42	43	45	46	47		
	38	39	40	42	43	44	46	47	48		

B) Reshape our kernel (flatten)

- Here is how it works:




Figure 2: Reshaped version of the 2 kernels

- As you can see, each filter is flattened and then stacked together. Thus, for X filter, we will flatten and stack X filters together.

C) Matrix multiplication between reshaped input and kernel

- Now, we only need to perform a matrix multiplication.




Figure 3: Matrix multiplication

- At the end, we need to reshape our matrix back to an image.
- Be aware the `np.reshape()` method doesn't return the expected result here (elements in wrong order). A little bit of numpy gymnastic solves the problem.

Implementation

- The most difficult part of `im2col` is to **transform our input image into a matrix**.
- To do so, we will use `np.tile()` and `np.repeat()` methods from Numpy. Here is how they work:

```
>>> y = np.arange(5)
>>> y
array([0, 1, 2, 3, 4])
>>> np.tile(y, 2)
array([0, 1, 2, 3, 4, 0, 1, 2, 3, 4])
>>> np.repeat(y, 2)
array([0, 0, 1, 1, 2, 2, 3, 3, 4, 4])
```

- Now, let's explain the process visually.

Reminder:

- We want to perform a convolution between an (1,3,4,4) input image kernels of shape (2,3,2,2).
- For index i:

Create level 1 vector
(first channel)

0	0	1	1
---	---	---	---



Expand to other channels

0	0	1	1	0	0	1	1	0	0	1	1
---	---	---	---	---	---	---	---	---	---	---	---

Create a vector with an increase by 1 at each level

level 1				level 2			level 3		
0	0	0	0	1	1	1	2	2	2

+

	0	0	0	1	1	1	2	2	2
0	0	0	0						
0	0	0	0						
1	1	1	1						
1	1	1	1						
0	0	0	0						
0	0	0	0						
1	1	1	1						
1	1	1	1						
0	0	0	0						
0	0	0	0						
1	1	1	1						
1	1	1	1						

Add both vectors to create matrix of index i at every levels for every channels.

level 1

			(0,0)	(0,1)	(0,2)	(0,3)
		(0,0)	(0,1)	(0,2)	(0,3)	(1,3)
	(0,0)	(0,1)	(0,2)	(0,3)	(1,3)	(2,3)
(1,0)	(1,1)	(1,2)	(1,3)	(2,3)	(3,3)	
(2,0)	(2,1)	(2,2)	(2,3)	(3,3)		
(3,0)	(3,1)	(3,2)	(3,3)			

- For index j:

Create slide 1 vector
(first channel)


0	1	0	1
---	---	---	---



Expand to other channels

0	1	0	1	0	1	0	1	0	1	0	1
---	---	---	---	---	---	---	---	---	---	---	---

Create a vector with an increase by 1 at each slide.
(Keep same pattern at each level)



+

	0	1	2	0	1	2	0	1	2
0	0	1	2						
1	1	2	3						
0	0	1	2						
1	1	2	3						
0	0	1	2						
1	1	2	3						
0	0	1	2						
1	1	2	3						
0	0	1	2						
1	1	2	3						
0	0	1	2						
1	1	2	3						
0	0	1	2						
1	1	2	3						

Add both vectors to create matrix of index j at every levels for every channels.

level 1

			(0,0)	(0,1)	(0,2)	(0,3)
		(0,0)	(0,1)	(0,2)	(0,3)	(1,3)
	(0,0)	(0,1)	(0,2)	(0,3)	(1,3)	(2,3)
(0,0)	(0,1)	(0,2)	(0,3)	(1,3)	(2,3)	(3,3)
(1,0)	(1,1)	(1,2)	(1,3)			
(2,0)	(2,1)	(2,2)	(2,3)			
(3,0)	(3,1)	(3,2)	(3,3)			

- Now, we can transform our input image into a matrix.

Matrix of index i

0	0	0	1	1	1	2	2	2
0	0	0	1	1	1	2	2	2
1	1	1	2	2	2	3	3	3
1	1	1	2	2	2	3	3	3
0	0	0	1	1	1	2	2	2
0	0	0	1	1	1	2	2	2
1	1	1	2	2	2	3	3	3
1	1	1	2	2	2	3	3	3
0	0	0	1	1	1	2	2	2
0	0	0	1	1	1	2	2	2
1	1	1	2	2	2	3	3	3
1	1	1	2	2	2	3	3	3

Matrix of index j

0	1	2	0	1	2	0	1	2
0	1	2	0	1	2	0	1	2
1	2	3	1	2	3	1	2	3
1	2	3	1	2	3	1	2	3
0	1	2	0	1	2	0	1	2
0	1	2	0	1	2	0	1	2
1	2	3	1	2	3	1	2	3
1	2	3	1	2	3	1	2	3
0	1	2	0	1	2	0	1	2
0	1	2	0	1	2	0	1	2
1	2	3	1	2	3	1	2	3
1	2	3	1	2	3	1	2	3

Here is the code to implement **im2col**:


```

1  def get_indices(X_shape, HF, WF, stride, pad):
2      """
3          Returns index matrices in order to transform our input image into a matrix.
4
5          Parameters:
6          -X_shape: Input image shape.
7          -HF: filter height.
8          -WF: filter width.
9          -stride: stride value.
10         -pad: padding value.
11
12         Returns:
13         -i: matrix of index i.
14         -j: matrix of index j.
15         -d: matrix of index d.
16             (Use to mark delimitation for each channel
17             during multi-dimensional arrays indexing).
18         """
19         # get input size
20         m, n_C, n_H, n_W = X_shape
21
22         # get output size
23         out_h = int((n_H + 2 * pad - HF) / stride) + 1
24         out_w = int((n_W + 2 * pad - WF) / stride) + 1
25
26         # ----Compute matrix of index i----
27
28         # Level 1 vector.
29         level1 = np.repeat(np.arange(HF), WF)
30         # Duplicate for the other channels.
31         level1 = np.tile(level1, n_C)
32         # Create a vector with an increase by 1 at each level.
33         everyLevels = stride * np.repeat(np.arange(out_h), out_w)
34         # Create matrix of index i at every levels for each channel.
35         i = level1.reshape(-1, 1) + everyLevels.reshape(1, -1)
36
37         # ----Compute matrix of index j----
38
39         # Slide 1 vector.
40         slide1 = np.tile(np.arange(WF), HF)
41         # Duplicate for the other channels.
42         slide1 = np.tile(slide1, n_C)
43         # Create a vector with an increase by 1 at each slide.
44         everySlides = stride * np.tile(np.arange(out_w), out_h)
45         # Create matrix of index j at every slides for each channel.
46         j = slide1.reshape(-1, 1) + everySlides.reshape(1, -1)
47
48         # ----Compute matrix of index d----
49
50         # This is to mark delimitation for each channel
51         # during multi-dimensional arrays indexing.
52         d = np.repeat(np.arange(n_C), HF * WF).reshape(-1, 1)
53
54         return i, j, d
55
56  def im2col(X, HF, WF, stride, pad):
57      """
58          Transforms our input image into a matrix.
59
60          Parameters:
61          - X: input image.
62          - HF: filter height.
63          - WF: filter width.
64          - stride: stride value.
65          - pad: padding value.
66
67          Returns:
68          -cols: output matrix.
69      """
70      # Padding
71      X_padded = np.pad(X, ((0,0), (0,0), (pad, pad), (pad, pad)), mode='constant')
72      i, j, d = get_indices(X.shape, HF, WF, stride, pad)
73      # Multi-dimensional arrays indexing

```

```

73 # multi-dimensional arrays indexing.
74 cols = X_padded[:, d, i, j]
75 cols = np.concatenate(cols, axis=-1)
76 return cols

```

2) Pooling layer

```

78 def forward(self, X):


```

- We can make the average pooling operation faster by using **im2col** method. Performs a forward convolution.
- Be aware that the `np.reshape()` method doesn't return the expected result here (elements in wrong order). **np.nditer** or of numpy gymnastic solves the problem.

```

83 - X : Last conv layer of shape (m, n_C_prev, n_H_prev, n_W_prev).

```



Here is the implementation code:

```

1 def forward(self, X):
2     """
3         Apply average pooling.
4
5         Parameters:
6         - X: Output of activation function.
7
8         Returns:
9         - A_pool: X after average pooling layer.
10    """
11    self.cache = X
12
13    m, n_C_prev, n_H_prev, n_W_prev = X.shape
14    n_C = n_C_prev
15    n_H = int((n_H_prev + 2 * self.p - self.f) / self.s) + 1
16    n_W = int((n_W_prev + 2 * self.p - self.f) / self.s) + 1
17
18    X_col = im2col(X, self.f, self.f, self.s, self.p)
19    X_col = X_col.reshape(n_C, X_col.shape[0]//n_C, -1)
20    A_pool = np.mean(X_col, axis=1)
21    # Reshape A_pool properly.
22    A_pool = np.array(np.hsplit(A_pool, m))
23    A_pool = A_pool.reshape(m, n_C, n_H, n_W)
24
25    return A_pool

```

II) Backward propagation

- This part will be tougher than the previous one.
- However, if you have read the [previous post](https://hackmd.io/@bouteille/ByusmjZc8) (https://hackmd.io/@bouteille/ByusmjZc8), about the naive implementation of Convolutional Neural network using Numpy, it should be fine.
- Along the way, you will often encounter a “**Be aware**” sentence about reshaping. I strongly advise you to run and play with `unit_tests.py` to understand why these numpy gymnastics were required.

1) Convolutional layer

Reminder:

- We performed a convolution between (1,3,4,4) input image and kernels of shape (2,3,2,2) which output an (2,3,3) image.
- During the backward pass, the (2,3,3) image contains the error/gradient (“**dout**”) which needs to be back-propagated to the:
 - (1,3,4,4) input image (layer).
 - (2,3,2,2) kernels.

★ Layer gradient: Intuition


- The formula to compute the layer gradient is:

$$\frac{\partial L}{\partial I} = \text{Conv}(K, \frac{\partial L}{\partial O})$$


- $\frac{\partial L}{\partial I}$: Input gradient.
- K : Kernels.
- $\frac{\partial L}{\partial O}$: Output gradient.
- Conv : Convolution operation.
- To do so, we will proceed as follow:
 - **A.** Reshape `dout` ($\frac{\partial L}{\partial O}$).
 - **B.** Reshape kernels `w` into single matrix `w_col`.
 - **C.** Perform matrix multiplication between reshaped `dout` and kernel.
 - **D.** Reshape back to image (`col2im`).
- We are going to see how it works intuitively and then how to implement it using Numpy.

A) Reshape dout

- During backward propagation, the output of the forward convolution contains the error that needs to be back-propagated.




B) Reshape w into w_col




C) Perform matrix multiplication between reshaped dout and w_col

- In order to perform the matrix multiplication, we need to transpose w_col .
- We will denote the output as dx_col .




- Notice that we are in fact, broadcasting the error in $dout$ to each kernel as we did in the naive implementation.



D) Reshape back to image (col2im).

- Here, **col2im** is more than a simple backward operation of **im2col**. Indeed, we have to take care of cases where errors will overlap with others.
- As we can see in the previous gif, the (14,14,6) image has overlapping window. We need to reproduce the same effect !



★ Layer gradient: Implementation

- The most difficult part of **col2im** is to reshape our matrix back to an image because it requires us to take care of the overlapping gradient.
- An efficient and elegant way to do so is to use the **np.add.at** (<https://numpy.org/doc/stable/reference/generated/numpy.ufunc.at.html>) method from Numpy. Here is a short example of how it works:

```
1 >>> indices = [  
2             [0,4,1], # rows.  
3             [3,2,4] # columns.  
4         ]  
5 >>> X = np.zeros((5,6))  
6 >>> np.add.at(X, indices, 1)  
7 >>> X  
8 array([[0, 0, 0, 1, 0, 0],  
9        [0, 0, 0, 0, 1, 0],  
10       [0, 0, 0, 0, 0, 0],  
11       [0, 0, 0, 0, 0, 0],  
12       [0, 0, 1, 0, 0, 0]])
```

- We will proceed as follow:
 - Create a matrix filled with 0 of the same shape as input image (add padding if needed).
 - **x_padded** : (1,3,4,4) with pad=0.
 - Use **get_indices()** which returns index matrices, necessary to transform our input image into a matrix.
 - **i** : (12,9)
 - **j** : (12,9)
 - **d** : (12,1)
 - Retrieve **dx_col** batch dimension by splitting it **N** (number of images) times. For example, if you have **N** images, then:
 - **dx_col** : (12, 9) => (N, 12, 9)
 - Be aware that the **np.reshape()** method doesn't return the expected result here (elements in wrong order). A little bit of numpy gymnastic solves the problem.
 - Use **i,j,d** matrices as argument in **np.add.at** to reshape our matrix back to input image.
 - Refer to step **D) Reshape back to image (col2im)** for **np.add.at** method visualization.

- Remove padding from new image if needed.

○ Kernel gradient: Intuition


- The formula to compute the kernel gradient is:

$$\frac{\partial L}{\partial K} = \text{Conv}(I, \frac{\partial L}{\partial O})$$

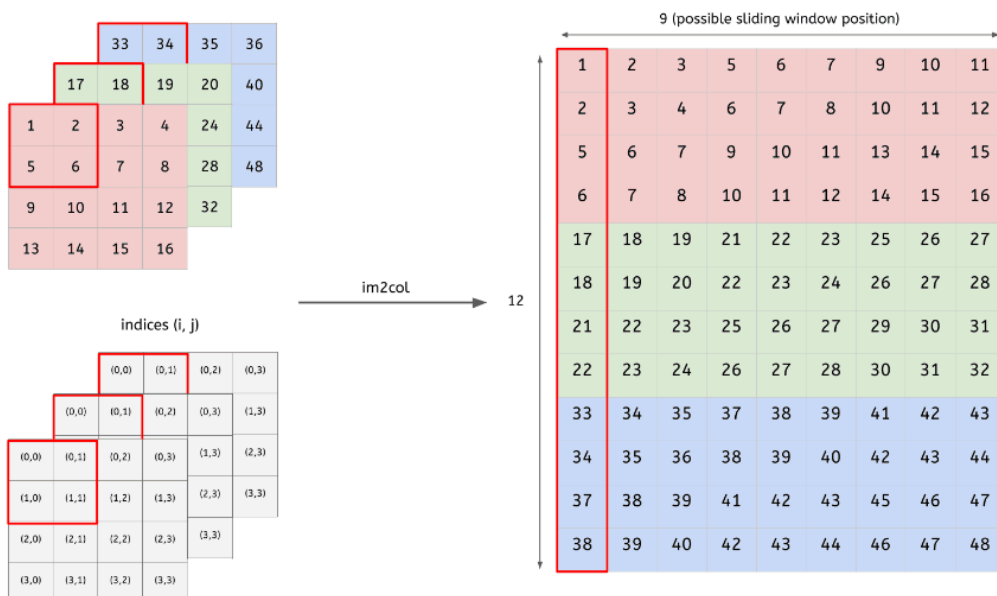
- $\frac{\partial L}{\partial K}$: Kernels gradient.
- I : Input image.
- $\frac{\partial L}{\partial O}$: Output gradient.
- Conv : Convolution operation.
- To do so, we will:
 - **A.** Reshape dout ($\frac{\partial L}{\partial O}$).
 - **B.** Apply **im2col** on x to get x_{col} .
 - **C.** Perform matrix multiplication between reshaped dout and x_{col} to get dw_col .
 - **D.** Reshape dw_col back to dw .
- We are going to see how it works intuitively and then how to implement it using Numpy.

A) Reshape dout

- Be aware that the `np.reshape()` method doesn't return the expect result here (elements in wrong order). A little bit of numpy gymnastic solves the problem.




B) Apply im2col on X to get X_col




C) Perform matrix multiplication between reshaped dout and X_col to get dw_col

- In order to perform the matrix multiplication, we need to transpose X_col .




- Notice that we are in fact, broadcasting the error in $dout$ to each "slide" we did during the naive implementation forward propagation over the input.



D) Reshape dw_col back to dw

- We simply need to reshape dw_col back to its original kernel shape.



○ Kernel gradient: Implementation

- Nothing fancy here.

Here is the code to implement the layer and kernel gradient.

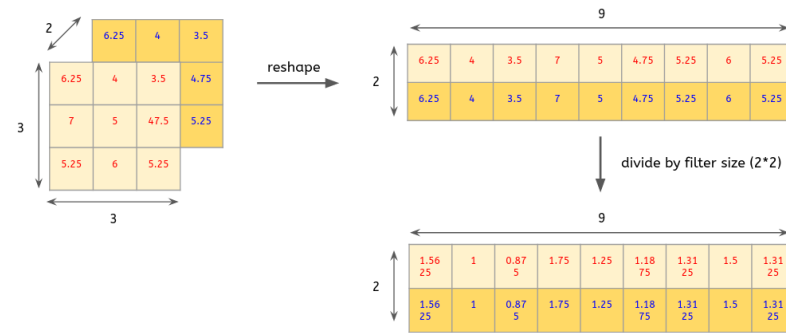
```

1  def col2im(dX_col, X_shape, HF, WF, stride, pad):
2      """
3          Transform our matrix back to the input image.
4
5          Parameters:
6          - dX_col: matrix with error.
7          - X_shape: input image shape.
8          - HF: filter height.
9          - WF: filter width.
10         - stride: stride value.
11         - pad: padding value.
12
13         Returns:
14         -x_padded: input image with error.
15     """
16     # Get input size
17     N, D, H, W = X_shape
18     # Add padding if needed.
19     H_padded, W_padded = H + 2 * pad, W + 2 * pad
20     X_padded = np.zeros((N, D, H_padded, W_padded))
21
22     # Index matrices, necessary to transform our input image into a matrix.
23     i, j, d = get_indices(X_shape, HF, WF, stride, pad)
24     # Retrieve batch dimension by splitting dX_col N times: (X, Y) => (N, X, Y)
25     dX_col_resaped = np.array(np.hsplit(dX_col, N))
26     # Reshape our matrix back to image.
27     # slice(None) is used to produce the [::] effect which means "for every elements"
28     np.add.at(X_padded, (slice(None), d, i, j), dX_col_resaped)
29     # Remove padding from new image if needed.
30     if pad == 0:
31         return X_padded
32     elif type(pad) is int:
33         return X_padded[pad:-pad, pad:-pad, :, :]
34
35 def backward(self, dout):
36     """
37         Distributes error from previous layer to convolutional layer and
38         compute error for the current convolutional layer.
39
40         Parameters:
41         - dout: error from previous layer.
42
43         Returns:
44         - dX: error of the current convolutional layer.
45         - self.W['grad']: weights gradient.
46         - self.b['grad']: bias gradient.
47     """
48     X, X_col, w_col = self.cache
49     m, _, _, _ = X.shape
50     # Compute bias gradient.
51     self.b['grad'] = np.sum(dout, axis=(0,2,3))
52     # Reshape dout properly.
53     dout = dout.reshape(dout.shape[0] * dout.shape[1], dout.shape[2] * dout.shape[3])
54     dout = np.array(np.vsplit(dout, m))
55     dout = np.concatenate(dout, axis=-1)
56     # Perform matrix multiplication between reshaped dout and w_col to get dX_col.
57     dX_col = w_col.T @ dout
58     # Perform matrix multiplication between reshaped dout and X_col to get dw_col.
59     dw_col = dout @ X_col.T
60     # Reshape back to image (col2im).
61     dX = col2im(dX_col, X.shape, self.f, self.f, self.s, self.p)
62     # Reshape dw_col into dw.
63     self.W['grad'] = dw_col.reshape((dw_col.shape[0], self.n_C, self.f, self.f))
64
65     return dX, self.W['grad'], self.b['grad']


```

2) Pooling layer

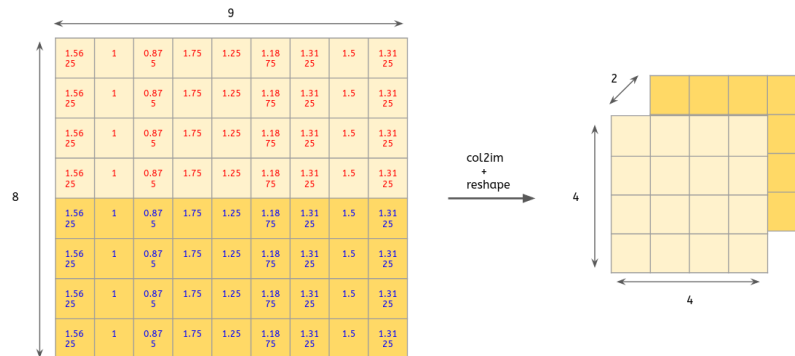
- We first have to reshape our filters and divide by the filter size.



- We then repeat each element "filter size" time.



- Finally, we apply `col2im`.
- Be aware that the `np.reshape()` method doesn't return the expected result here (elements in wrong order). A little bit of numpy gymnastic solves the problem.



Here is the implementation code:

```

1  def backward(self, dout):
2      """
3          Distributes error through pooling layer.
4
5          Parameters:
6          - dout: Previous layer with the error.
7
8          Returns:
9          - dX: Conv layer updated with error.
10     """
11     X = self.cache
12     m, n_C_prev, n_H_prev, n_W_prev = X.shape
13
14     n_C = n_C_prev
15     n_H = int((n_H_prev + 2 * self.p - self.f) / self.s) + 1
16     n_W = int((n_W_prev + 2 * self.p - self.f) / self.s) + 1
17
18     dout_flatten = dout.reshape(n_C, -1) / (self.f * self.f)
19     dX_col = np.repeat(dout_flatten, self.f*self.f, axis=0)
20     dX = col2im(dX_col, X.shape, self.f, self.f, self.s, self.p)
21     # Reshape dX properly.
22     dX = dX.reshape(m, -1)
23     dX = np.array(np.hsplit(dX, n_C_prev))
24     dX = dX.reshape(m, n_C_prev, n_H_prev, n_W_prev)
25     return dX

```

III) Performance of fast implementation

- The naive implementation (<https://github.com/3outeille/CNNumpy/tree/master/src/slow>) takes around **4 hours for 1 epoch** where the fast implementation (<https://github.com/3outeille/CNNumpy/tree/master/src/fast>) takes only **6 min for 1 epoch**.
- For your information, **with the same architecture using Pytorch**, it will take around **1 min for 1 epoch**.